

Instruction Prefetching Algorithms Comparison on PIN



(Amanda) Yilan Zhu
Northeastern University
EECE 7352 Computer Architecture

Introduction

- Instruction cache misses are crucial when evaluating the performance of microprocessors.
- Different strategies: Increase cache size/associativity; adding a victim cache/branch prediction...
- Cache prefetching: 2 tasks:
 - To “guess” which line will be soon accessed
 - Initiate the prefetch “just in time”
- If done correctly, prefetching could remove all cache misses.

Project Goals

- Simulate 3 cache prefetching algorithms:
 - Next-line prefetching
 - Target-line prefetching
 - Wrong-path prefetching
- Benchmark
 - Barnes.c - Simulates the gravitational forces acting on a star cluster using the Barnes-Hut n-body algorithm.
 - FindPi.c – Find the value of Pi by the Monte-Carlo method.
- Different Cache Set Associativity
 - Direct mapping/ 2-way/ 4-way

Tools

- Intel Pin 3.0
 - Tool for the “dynamic instrumentation of programs”,
 - Customize you own pintool to “simulate an instruction cache.”
 - Ran on COE nano (x86_64)
- Pintool : icache
 - Instruction cache simulator (icache.cpp)
 - Available in /pin/sources/tools/Memory
 - Modified to have prefetching abilities
 - Modifications made in both instrumentation and analysis routines

Prefetching Algorithms - 1

- Next-line prefetching
 - Prefetch next line if not in cache and when current line is in certain distance of next line.

In Analysis routine: executed at insertion points

```
VOID LoadSingle(ADDRINT addr, ADDRINT next_addr, UINT32 size, UINT32 instId)
{
    const BOOL Hit = ill->AccessSingleLine(next_addr, CACHE_BASE::ACCESS_TYPE_LOAD);
    if (Hit == 0)
    {
        if ((next_addr - addr) < 2 * size){
            const BOOL MyHit = ill->AccessSingleLine(addr, CACHE_BASE::ACCESS_TYPE_LOAD);
            const COUNTER counter = MyHit ? COUNTER_HIT : COUNTER_MISS;
            profile[instId][counter]++;
        }
    }
}
```

When the next line is not in cache (miss), it will be prefetched when the distance between the current and next address is smaller than 2 cache line sizes.

Prefetching Algorithms - 2

- Target-Line prefetching
 - Uses a maintained target table (current line – successor)
 - If successor line of current address in table is not in cache, prefetch.

```
VOID LoadSingle(ADDRINT next_addr, ADDRINT addr, UINT32 instId)
```

In Analysis routine: executed at insertion points

```
{
    unsigned long int Target_table[100][2] = {0};
    BOOL target;
    int I, n = 0;
    for (int i = 0; i < 100; i++)
    {
        target = (addr != Target_table[i][0]);
        if (target == 0)
            I = i;
    }

    if (target)
    {
        Target_table[n][0] = addr;
        Target_table[n][1] = next_addr;
        const BOOL Hit1 = ill->AccessSingleLine(addr, CACHE_BASE::ACCESS_TYPE_LOAD);
        const COUNTER counter = Hit1 ? COUNTER_HIT : COUNTER_MISS;
        profile[instId][counter]++;
        n++;
    }

    else
    {
        const BOOL Hit2 = ill->AccessSingleLine(Target_table[I][1], CACHE_BASE::ACCESS_TYPE_LOAD);
        const COUNTER counter = Hit2 ? COUNTER_HIT : COUNTER_MISS;
        profile[instId][counter]++;
    }
}
```

Creates a fixed size target table with current address and successor line.

This will be passed to the instrumentation analysis function.

Prefetching Algorithms - 3

- Wrong-path prefetching
 - Combination of next-line and target prefetching.
 - Not-taken branch -> prefetch cache line with fallthrough address
 - Taken branch -> prefetch branch target line
 - Perform better than correct-path only schemes when there exists a large disparity between the CPU cycle time and the memory speed.
 - This is because other algorithms try to prefetch down target paths which will be executed almost immediately, and if memory is slow the prefetch may not be initiated soon enough.

Prefetching Algorithms - 3

Instrumentation: A mechanism that decides where and what code is inserted.

```
VOID Instruction(INS ins, void * v)
```

```
{  
    // map sparse INS addresses to dense IDs  
    const ADDRINT iaddr = INS_Address(ins);  
    const UINT32 instId = profile.Map(iaddr);
```

```
  
    const UINT32 size = INS_Size(ins);  
    const BOOL single = (size <= 4);
```


```
  
    if (KnobTrackInsts) {  
        if (single)  
        {  
            if(INS_IsBranch(ins) && INS_HasFallThrough(ins)){ //if is taken branch  
                BOOL my_Hit1 = ill->AccessSingleLine(INS_DirectBranchOrCallTargetAddress(ins), CACHE_BASE::ACCESS_TYPE_LOAD);  
                const COUNTER counter = my_Hit1 ? COUNTER_HIT : COUNTER_MISS;  
                profile[instId][counter]++;  
            }  
            else{  
                BOOL my_Hit2 = ill->AccessSingleLine(INS_NextAddress(ins), CACHE_BASE::ACCESS_TYPE_LOAD);  
                const COUNTER counter = my_Hit2 ? COUNTER_HIT : COUNTER_MISS;  
                profile[instId][counter]++;  
            }  
        }  
        else {  
            INS_InsertPredicatedCall(ins, IPOINT_BEFORE, (AFUNPTR) LoadMulti,  
                                     IARG_UINT32, iaddr,  
                                     IARG_UINT32, size,  
                                     IARG_UINT32, instId,  
                                     IARG_END);
```

If is taken branch, fetch the target address;
if not taken branch, fetch next “natural” address

Experiment Design

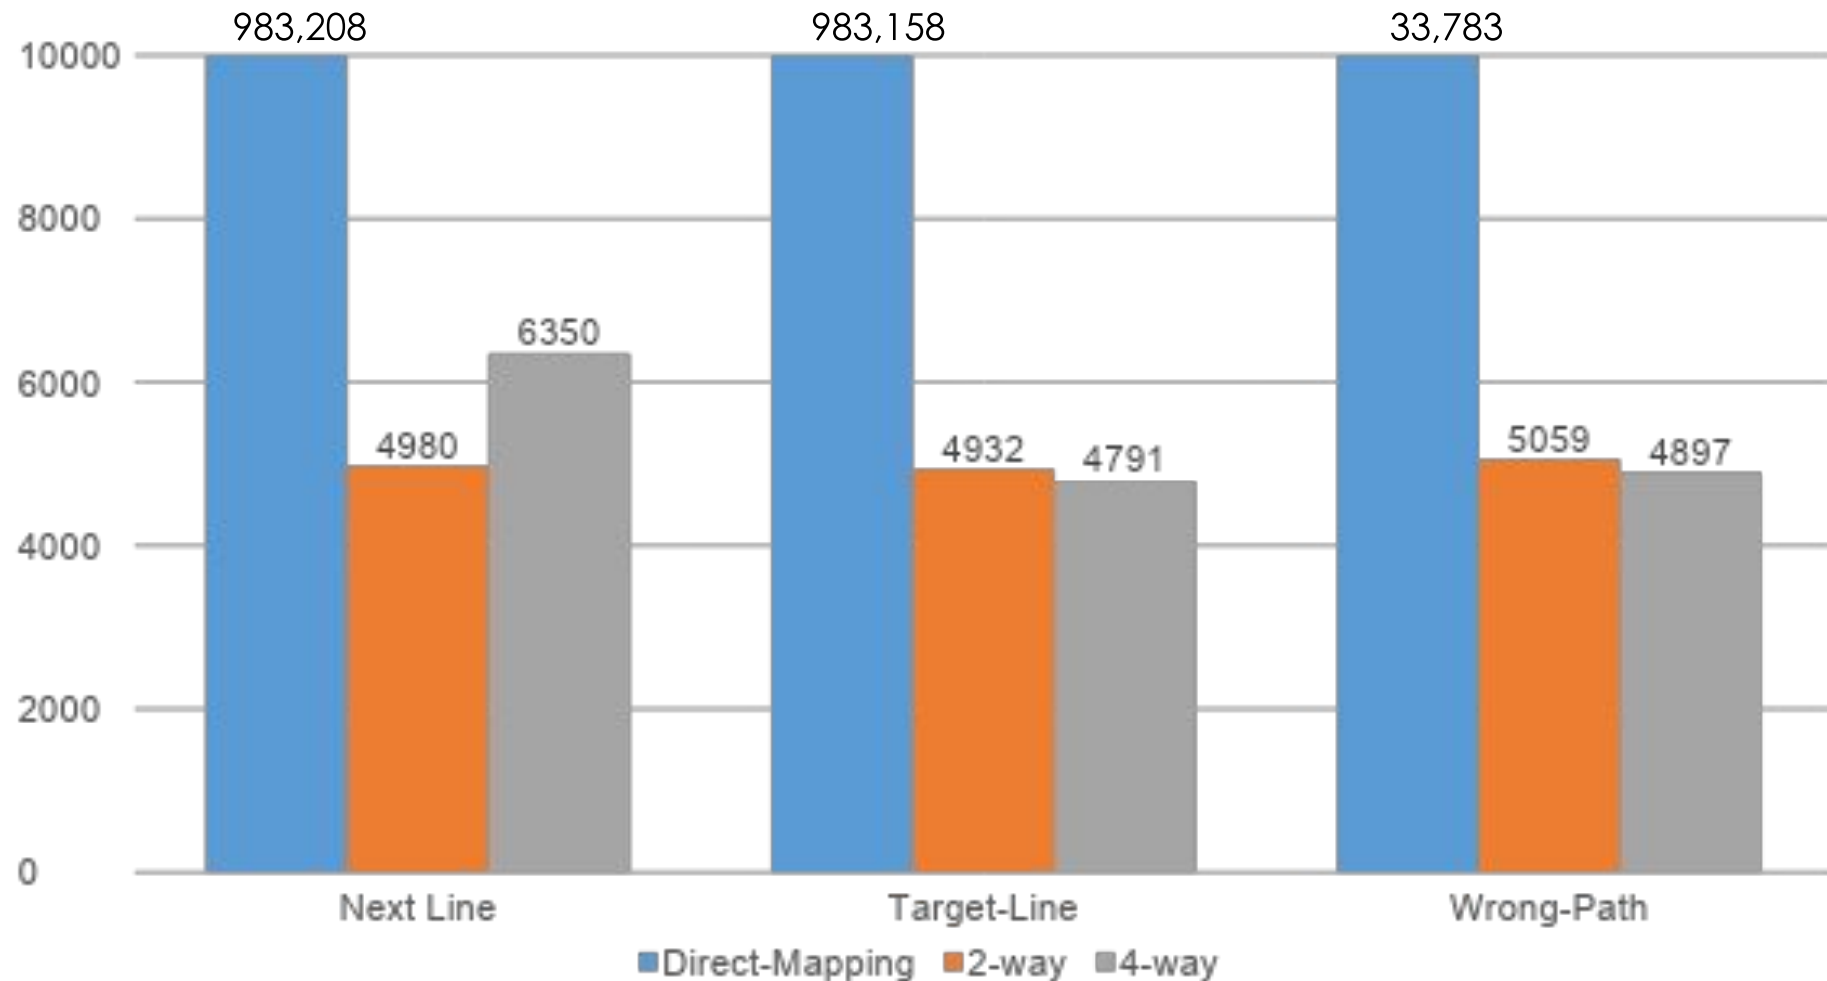
- Next-Line prefetching and Target-line prefetching done on both analysis and instrumentation routine.
- Wrong-path only done on instrumentation.
- Direct-mapping/ 2-way / 4-way set associative.
- Benchmarks: FindPi.c / Barnes

Results




Results

Instruction misses in Barnes - Instrumentation

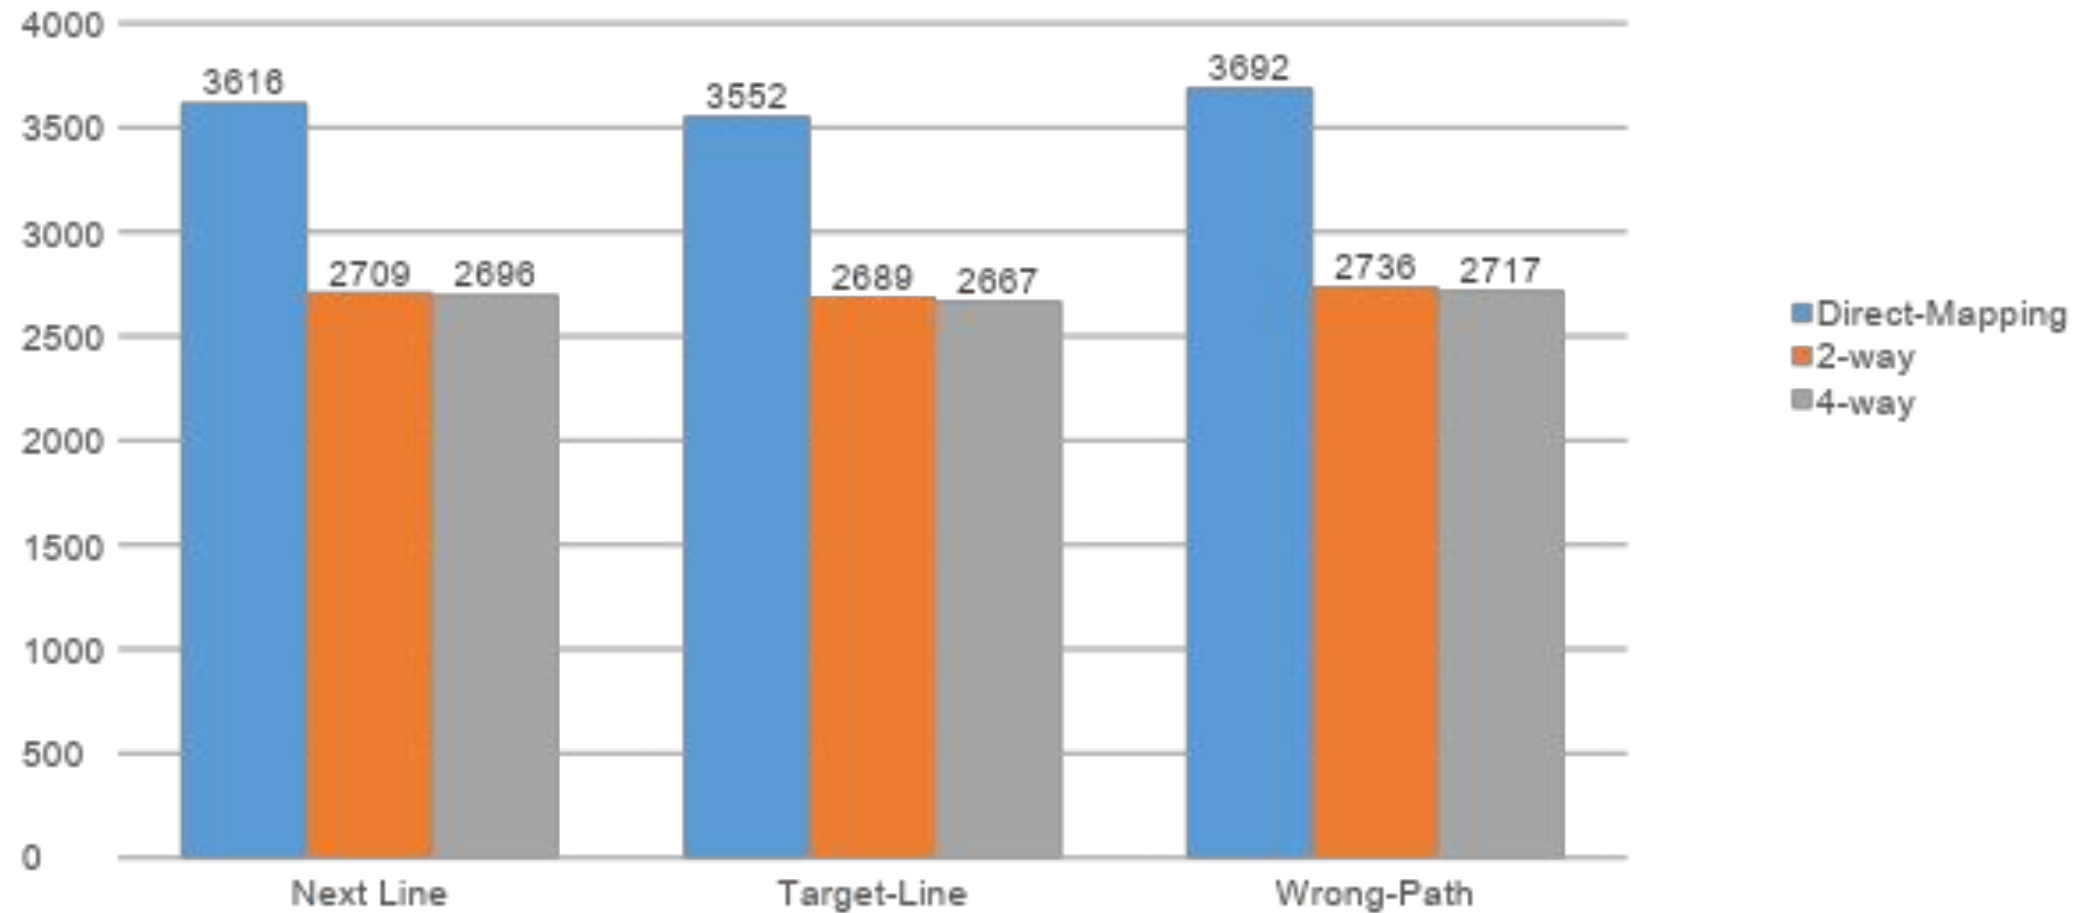


Results



Results

Instruction misses in FindPi - Instrumentation



Conclusions

- Set-associative improve cache hits significantly.
- Wrong Path prefetching, slightly improves cache hits
- Pintool modification should be done in Analysis.